

Objective-Dependent Active Learning and Calibrated Uncertainty for Sample-Efficient Discovery of Rare-Earth Extractants: A Benchmark on 1,202 Experimental Distribution Coefficients with Prospective Validation

Krishna Harish*

Lawrence E. Elkins High School, Missouri City, Texas, USA

E-mail: krishnahaarish2009@gmail.com

Abstract

Measuring distribution coefficients ($\log D$) for rare-earth solvent extraction is slow and expensive, making sample-efficient machine learning attractive, but the practical questions of *which* acquisition strategy to use, whether the attendant uncertainty estimates can be trusted, and how well a model actually generalises to new chemistry are rarely answered together on real data. We benchmark active learning (AL) on 1,202 experimental $\log D$ measurements for lanthanide extraction (93 ligands, 14 lanthanides, 45 literature sources), with prospective validation on four independently synthesised ligands. Five findings emerge, each on real data with multi-seed statistics. (i) A plain random forest reaches $R^2 = 0.94$ /RMSE= 0.33 on the published row-level validation split, outperforming the previously reported deep network ($R^2 = 0.85$); however, under leakage-controlled splits that forbid a ligand, source, or structural family from appearing in both train and test, the estimate collapses (grouped-CV $R^2 = 0.20$ to -0.03),

so the in-distribution number is strongly optimistic and realistic deployability is far lower. (ii) The optimal AL acquisition is *objective-dependent*: exploitation recovers all top row-level extractants (top-10 recall = 1.0) but yields a biased, poorly predictive model, whereas exploration/random gives the best predictive model but finds few of the best; no strategy dominates both, and this survives a chemically realistic ligand-batch acquisition unit and reproduces on an independent 4,200-compound $\log D$ benchmark. (iii) The discovery ranking of strategies *inverts* with the definition of a “top” system: exploitation dominates for condition-specific rows, but broad sampling already recovers most top ligands. (iv) For *prediction*, AL provides no advantage over random selection. (v) Conformal calibration is essential for trustworthy uncertainty (an eight-fold in-distribution error reduction) but does not improve AL acquisition, which depends only on the uncertainty ranking; under distribution shift, global conformal leaves subgroup imbalance that group-conditional (Mondrian) calibration reduces, at the cost of reordering the ranking. We compare five model/uncertainty estimators, add early-recognition discovery metrics, and distil the results into a practical decision guide, releasing all code and data.

1 Introduction

Separating rare-earth and f -elements by solvent extraction underpins critical-metal supply and nuclear-fuel-cycle chemistry, and the key performance descriptor, the distribution coefficient $\log D$, is measured one ligand/metal/condition at a time by liquid–liquid assay.¹ Each measurement is laborious, so a model that reaches useful accuracy, or surfaces the best extractants, from *few* measurements has direct practical value. Machine learning for $\log D$ and distribution-coefficient prediction in rare-earth extraction has advanced quickly,^{2–5} and active learning (AL)^{6,7} promises to choose the most informative experiments next; pool-based AL with surrogate models has accelerated molecular discovery and high-throughput screening,^{8–10} and batch acquisition is now standard when several experiments are run per

round.¹¹

Three practitioner-facing questions are nonetheless usually left unanswered on real extractant data. First, *which* acquisition strategy to use when the goal may be accurate prediction *or* discovery of top performers: the exploration –exploitation trade-off is well studied in Bayesian optimisation,^{12–14} and sequential-learning benchmarks show the best scheme is goal-dependent and can even underperform random,⁹ but this is rarely quantified for f -element separation. Second, *whether* the uncertainty estimates AL relies on are calibrated enough to trust: conformal prediction gives distribution-free guarantees^{15–17} and is increasingly used in cheminformatics,^{18–20} yet valid average coverage can still hide failures in the high-log D region that matters most for discovery. Third, *how well* a model generalises beyond its training chemotypes: random cross-validation is known to overstate prospective performance,²¹ and scaffold or cluster splits expose generalization gaps that random splits hide.^{22–24} We address all three on a real benchmark of 1,202 experimental log D values with prospective validation, and report findings that are useful precisely because several run counter to the common framing of AL as a uniformly beneficial drop-in.

2 Data and Methods

Dataset and recovered group structure. We use the 1,202 experimental lanthanide-extraction log D measurements compiled by Liu et al.² (1,085 training + 117 validation), with their feature representation: Morgan/ECFP bits and RDKit physicochemical descriptors for the ligand, experimental conditions (ligand and acid concentration, diluent identity/ratio, temperature), and lanthanide-identity descriptors (2,291 features per row). The released matrix contains no SMILES, but the group structure needed for rigorous evaluation is recoverable from it: identical ECFP bit-vectors identify the 93 unique ligands, the metal-descriptor block identifies the 14 lanthanides, and the integer citation index identifies the 45 publication sources (median 14 measurements and 9 distinct lanthanides per ligand). Their

four subsequently synthesised ligands (56 ligand $\times$ metal measurements) serve as a strictly prospective external test.

Leakage-controlled splits. Because random row-level cross-validation can place measurements of the same ligand or source in both train and test, we evaluate generalisation under four regimes with identical models: plain 5-fold over rows, and group-disjoint 5-fold (repeated with reshuffled group assignments) that forbid a shared *ligand*, *source*, or structural *family*. Exact Bemis–Murcko scaffolds require SMILES that are not distributed, so we use ECFP-Tanimoto agglomerative clustering (20 families at similarity ≥ 0.35) as an explicit scaffold-split proxy.

Surrogate, models, and uncertainty. The base model is a random-forest regressor²⁵ (scikit-learn²⁶); epistemic uncertainty is the standard deviation across trees, an ensemble estimator in the spirit of deep ensembles.^{27–29} To probe whether the conclusions depend on this choice we additionally benchmark a quantile regression forest, a Gaussian process (on PCA-reduced features), gradient boosting, and a Monte-Carlo-dropout Bayesian neural network,^{30–32} scored on predictive accuracy, calibration, and uncertainty-ranking quality. We apply split-conformal calibration^{15–17} (nonconformity $|y - \hat{y}|/\hat{\sigma}$ on a held-out calibration split), and, for subgroup coverage, group-conditional (Mondrian) conformal prediction.²⁰

Active-learning protocol. From a small random seed set we iteratively acquire batches and retrain. We use two acquisition units: the original row unit (fixed 80/20 pool/test partition) and a chemically realistic *ligand-batch* unit, in which selecting a ligand reveals all of its measurements at once and the test set is a disjoint set of held-out ligands. Acquisition strategies: *random*; *greedy* (exploit, maximise predicted $\log D$); *uncertainty* (explore, maximise $\hat{\sigma}$); *UCB* ($\hat{\mu} + 1.5\hat{\sigma}$); *expected improvement* (EI);^{13,14} and a *hybrid* portfolio (half exploit, half explore). We track, versus number of measurements, the held-out predictive R^2 and discovery metrics: top-10 recall under three definitions of a “top system” (condition-specific row, ligand–metal pair, unique ligand), the enrichment factor,^{33,34} simple regret, and hit-rate above a practical threshold ($\log D \geq 1$, i.e. $D \geq 10$). All curves are averaged over

six random seeds.

3 Results and Discussion

A strong but optimistic simple baseline. A random forest on the published features attains $R^2 = 0.942$, RMSE= 0.328 on the published validation split, exceeding the reported deep network ($R^2 = 0.85$, RMSE= 0.53).² For this tabular, condition-rich representation, ensemble trees are a stronger and far cheaper baseline. This comparison is on the same row-level split as the original, and is therefore an apples-to-apples model comparison, not an estimate of deployability.

Realistic generalisation under leakage-controlled splits (Figure 1, Table 1).

The row-level number is strongly optimistic. Holding out entire ligands, sources, or structural families collapses the estimate from $R^2 = 0.88$ (random rows) to $R^2 = 0.20$ (ligand-disjoint), -0.03 (source-disjoint), and -0.24 (family-disjoint), i.e. to at or below the performance of predicting the mean. The prospective test on four genuinely new ligands gives $R^2 = 0.70$, intermediate between the two extremes because those ligands are close analogues of training chemotypes. The practical implication is that in-distribution R^2 on this benchmark should not be read as deployability; realistic generalisation to new ligand chemistry is markedly lower, consistent with the random-split optimism documented across molecular ML.^{21,23}

Table 1: Random-forest generalisation under increasingly stringent, leakage-controlled splits. Random row-level cross-validation is strongly optimistic; grouped splits that forbid a ligand, source, or structural family from appearing in both train and test collapse the estimate toward the mean. Mean $\pm$ s.d. over 4 repeats of 5-fold grouped CV.

Split	R^2	RMSE	MAE
Random (row-level)	0.88 $\pm$ 0.03	0.46	0.28
Ligand-disjoint	0.20 $\pm$ 0.31	1.14	0.87
Source-disjoint	-0.03 $\pm$ 0.80	1.22	0.98
Ligand-family (scaffold proxy)	-0.24 $\pm$ 0.50	1.31	1.04
Prospective (4 new ligands)	0.70	0.63	0.49




Figure 1: Random-forest test R^2 under increasingly stringent, leakage-controlled splits. Random row-level CV is strongly optimistic; grouped splits that forbid a shared ligand, source, or structural family collapse the estimate toward the mean. Error bars are ± 1 s.d. over repeated grouped folds.

The optimal acquisition is objective-dependent (Figure 2). No single strategy is best. Exploitation (greedy, UCB) recovers *all* top row-level extractants (top-10 recall = 1.00) but biases the training set toward high-log D systems, giving a poor global model ($R^2 = 0.63$ – 0.70). Exploration and random give the best predictive models ($R^2 = 0.85$ – 0.87) but poor row-level design recall (0.38–0.50). EI and the hybrid portfolio sit on the Pareto front but dominate neither extreme. Practitioners must therefore choose the acquisition by the goal: exploit to *find* extractants, explore to *model* them, EI/hybrid to balance. This matches sequential-learning benchmarks in which the best scheme is goal-dependent and mismatched schemes underperform random.^{9,12}

The dichotomy survives a realistic ligand-batch unit (Figure 3). In the laboratory one selects a *ligand* and then measures it against many lanthanides and conditions, so we repeat the benchmark with a ligand-batch acquisition unit and a ligand-disjoint test set. The design-side dichotomy is clear: exploitation reaches the highest top-10 ligand recall (1.00) while random/exploration lags on recall (0.83). On the prediction side, all predictive R^2 are low under this deliberately hard ligand-disjoint test (consistent with the split analysis above), with random giving the best model ($R^2 = 0.30$) but by a small margin




Figure 2: Active learning on 1,202 real $\log D$ measurements (row unit). Left: predictive R^2 vs. number of measurements (exploration/random best). Right: top-10 recall (exploitation best). No strategy wins both; EI and the hybrid portfolio trace the Pareto front. Bands are ± 1 s.d. over six seeds.

over exploitation (0.29). The objective-dependent ordering is therefore preserved under the operational decision unit, not an artefact of scoring isolated rows, though the prediction-side gap narrows once the test set forbids train–test ligand overlap.

What counts as a “top” system changes the ranking (Figure 4, Table 2). Discovery value depends on whether a top system is a condition-specific row, a ligand–metal pair, or a unique ligand, and the strategy ranking *inverts* across these definitions. Exploitation dominates at the row level (greedy row-recall 1.00 vs random 0.45), because it targets the exact high- $\log D$ conditions. But at the ligand level the ordering flips: broad sampling already recovers most top ligands (random ligand-recall 0.95 vs greedy 0.82), because each ligand appears under many conditions and is therefore easy to hit without targeting. Reporting a single “top-10 recall” without stating the unit can thus reverse the practical conclusion; we report all three.

The dichotomy is general (Figure 5). On an independent dataset, the 4,200 octanol–water $\log D$ values from MoleculeNet’s Lipophilicity set²² with a different feature representation, the same anti-correlated pattern emerges: exploration and random give the best predictive models ($R^2 \approx 0.50$) but the lowest design recall (0.20–0.27), whereas exploitation




Figure 3: Ligand-batch acquisition (each acquired ligand reveals ~ 14 measurements), with a ligand-disjoint test set. Left: predictive R^2 ; right: top-10 ligand recall. The objective-dependent dichotomy persists under the realistic decision unit. Bands are ± 1 s.d. over six seeds.




Figure 4: Final top-10 recall under three definitions of a “top” system. Exploitation is best for condition-specific rows; broad/random sampling is competitive or better for unique ligands. The discovery conclusion depends on the definition.

Table 2: Discovery performance at the final budget under three definitions of a ‘top-10 system’ (condition-specific row, ligand–metal pair, unique ligand), plus simple regret (log units) and hit-rate above a practical threshold ($\log D \geq 1$, i.e. $D \geq 10$). The ranking of strategies inverts with the definition: exploitation dominates for rows, but broad sampling already recovers most top ligands. Means over six seeds.

Strategy	recall (row)	recall (lig–metal)	recall (ligand)	regret (log)	hit-rate ($\log D \geq 1$)
random	0.45	0.72	0.95	0.15	0.26
greedy	1.00	0.88	0.82	0.00	0.60
uncertainty	0.42	0.85	0.93	0.12	0.26
ucb	1.00	0.90	0.88	0.00	0.59
ei	0.93	0.90	0.97	0.02	0.47
hybrid	0.80	0.87	0.97	0.01	0.52

gives the best recall (0.53–0.57) but the worst predictive models ($R^2 = 0.21$ –0.28). The objective-dependence of the optimal acquisition is therefore a general property of pool-based AL, not an artefact of one dataset or feature set.




Figure 5: Generality. The identical prediction–design dichotomy on an independent benchmark (4,200 octanol–water $\log D$ values, Morgan-fingerprint features).

Active learning does not help prediction (a caution). At every budget, random selection matches or beats uncertainty sampling for predictive R^2 (e.g. 0.54 vs 0.40 at 100 measurements; both reach $R^2 = 0.85$ at ≈ 400). For building a predictive $\log D$ model, the experimental effort of AL is not repaid here, echoing UQ comparisons in which no acquisition uniformly helps.²⁹

Calibration matters for trust, not acquisition (Figure 6, Table 3). On the row-level split, raw ensemble uncertainty is badly miscalibrated (mean calibration error 0.103; the nominal-80% interval covers 90%), and split-conformal calibration reduces the error 8-fold to 0.013. Yet calibrating the uncertainty leaves AL performance unchanged, because acquisition depends only on the *ranking* of uncertainty, which a single global (monotone) conformal factor preserves. Under a realistic ligand-disjoint split the picture is richer: raw intervals now *under-cover* (0.72 at nominal 0.80), a single global conformal factor restores average coverage (0.90) but leaves the high-log D region unevenly covered, and group-conditional (Mondrian) conformal (0.87)²⁰ evens the subgroups. Crucially, Mondrian *reorders* the uncertainty ranking (Spearman $\rho = 0.85$ vs the global factor), so locally adaptive calibration *can* change AL acquisition even though a global factor cannot. Coverage on the prospective ligands is 0.88. Calibration is thus essential when the uncertainty is reported or acted upon,^{35,36} and, for acquisition, only non-monotone calibration matters, a distinction easy to conflate.




Figure 6: Subgroup coverage of nominal-80% intervals under a ligand-disjoint split. Raw intervals under-cover; a single global conformal factor restores average coverage but the high-log D region remains uneven; group-conditional (Mondrian) conformal balances the subgroups. Dashed line is the nominal target.

Broader model and uncertainty-estimator baselines (Table 4, Figure 7). To test whether the conclusions depend on the random forest, we compare five estimators on

Table 3: Uncertainty calibration under a realistic ligand-disjoint split (nominal coverage 80%). Raw ensemble intervals under-cover; a single global split-conformal factor restores average coverage but leaves subgroup imbalance, which group-conditional (Mondrian) conformal reduces. Mean over eight seeds. For reference, on the easier row-level split the raw mean calibration error is 0.103 and conformal reduces it to 0.013.

Method	global	low log D	mid log D	high log D
Raw	0.72	–	–	–
Global conformal	0.90	0.87	0.92	0.87
Mondrian conformal	0.87	0.96	0.82	0.88

a ligand-disjoint split: random forest, quantile regression forest, Gaussian process, gradient boosting, and an MC-dropout Bayesian neural network. No estimator uniformly dominates, consistent with prior UQ comparisons;^{28,29} gradient boosting gives the best uncertainty-ranking quality (Spearman between predicted σ and realised absolute error, $\rho = 0.28$). The dichotomy and the calibration-vs-acquisition distinction are properties of pool-based AL, not of a single surrogate.




Figure 7: Model/uncertainty-estimator comparison on a ligand-disjoint split: conformal 80% coverage and uncertainty-ranking quality (Spearman of predicted σ vs realised error). Error bars are ± 1 s.d. over four seeds.

Prospective validation. Trained on all 1,202 measurements and applied to the four independently synthesised ligands, the model attains $R^2 = 0.70$, RMSE= 0.639, MAE= 0.502 log units. Read together with the grouped-CV collapse above, this places realistic

Table 4: Model and uncertainty-estimator comparison on a ligand-disjoint split (train/calibration/test = 55/20/25 of ligands). Uncertainty-ranking quality is the Spearman correlation between predicted σ and realised absolute error; higher is better. Coverage is the empirical coverage of split-conformal 80% intervals. Mean over four seeds.

Model	R^2	conformal 80% cov.	rank quality $\rho(e , \sigma)$	enrich. factor
Random forest	0.15	0.91	0.22	7.5
Quantile RF	0.14	0.91	0.22	7.5
Gaussian process	-0.03	0.95	0.23	3.7
Gradient boosting	0.22	0.91	0.28	7.5
MC-dropout BNN	0.12	0.94	0.06	6.6

generalisation between the optimistic row-level number and the worst-case family-held-out number, and is, to our knowledge, among the few prospective tests reported for extractant-property models.

Practical guidance. (1) Start with an ensemble-tree baseline, but report a leakage-controlled (ligand/source/scaffold) split, not row-level CV, as the generalisation estimate. (2) Pick the acquisition by objective: exploit for discovery, explore/random for a predictive model, EI for balance. (3) State the unit of a “top” system; the discovery ranking depends on it. (4) Do not expect AL to reduce the measurements needed for accurate prediction. (5) Calibrate (e.g. conformally) whenever the uncertainty is reported or acted upon, and use group-conditional calibration when subgroup coverage matters; only non-monotone calibration changes acquisition.

4 Limitations

The objective-dependent dichotomy is shown on two independent datasets with different feature sets, but broader chemistries remain to be tested. Because SMILES are not distributed with the benchmark, our structural split uses ECFP-Tanimoto families as a scaffold-split proxy rather than exact Bemis–Murcko scaffolds. AL here is pool-based selection over measured candidates, not open-ended generative design, and prospective validation covers four

ligands. The conclusions concern $\log D$ regression and may differ for classification or multi-objective selectivity targets.

5 Conclusion

On real experimental extractant data with prospective validation, we show that a simple ensemble beats a published deep network but that row-level R^2 is strongly optimistic relative to leakage-controlled splits; that the best active-learning acquisition depends on whether the goal is prediction or discovery, with no strategy dominating both, and that this survives a realistic ligand-batch unit while the discovery ranking itself depends on how a “top” system is defined; that active learning offers no predictive-accuracy advantage over random selection; and that uncertainty calibration is essential for trust but changes acquisition only when it is non-monotone. These give practitioners an evidence-based decision guide and a reproducible benchmark for sample-efficient extractant discovery.

Data and Software Availability

All analysis code, recovered group labels, split definitions, and figure- and number-generating scripts are released at <https://github.com/krishnatheaverage/logd-active-learning-benchmark> and archived on Zenodo, reproducing every value herein on a single laptop. The underlying measurements are from ref. 2.

Acknowledgement

Computations were run on a single Apple M4 laptop. The author conceived and directed the study, curated the data, made the scientific and methodological decisions, and verified all results; in accordance with ACS policy, the author discloses that a generative AI assistant (Claude, Anthropic) was used under the author’s direction for code implementation,

experiment execution and analysis, and manuscript drafting. The author is accountable for the content.

References

- (1) Panak, P. J.; Geist, A. Complexation and Extraction of Trivalent Actinides and Lanthanides by Triazinylpyridine π -Donor Ligands. *Chemical Reviews* **2013**, *113*, 1199–1236.
- (2) Liu, T.; Johnson, K. R.; Jansone-Popova, S.; Jiang, D.-e. Advancing Rare-Earth Separation by Machine Learning. *JACS Au* **2022**, *2*, 1428–1434.
- (3) Udawattha, D. S.; Alam, S. Predicting the distribution coefficient in the solvent extraction of rare earth elements. *Separation and Purification Technology* **2025**, *377*, 134382.
- (4) Chaube, S.; Goverapet Srinivasan, S.; Rai, B. Applied machine learning for predicting the lanthanide-ligand binding affinities. *Scientific Reports* **2020**, *10*, 14322.
- (5) Edaugal, J. P.; Zhang, D.; Liu, D.; Glezakou, V.-A.; Sun, N. Solvent Screening for Separation Processes Using Machine Learning and High-Throughput Technologies. *Chem & Bio Engineering* **2025**,
- (6) Settles, B. *Active Learning Literature Survey*; 2009.
- (7) Reker, D. Practical considerations for active machine learning in drug discovery. *Drug Discovery Today: Technologies* **2020**, *32–33*, 73–79.
- (8) Graff, D. E.; Shakhnovich, E. I.; Coley, C. W. Accelerating high-throughput virtual screening through molecular pool-based active learning. *Chemical Science* **2021**, *12*, 7866–7881.
- (9) Rohr, B.; Stein, H. S.; Guevarra, D.; Wang, Y.; Haber, J. A.; Aykol, M.; Suram, S. K.; Gregoire, J. M. Benchmarking the acceleration of materials discovery by sequential learning. *Chemical Science* **2020**, *11*, 2696–2706.

- (10) Bayley, O.; Savino, E.; Slattery, A.; Noël, T. Autonomous chemistry: Navigating self-driving labs in chemical and material sciences. *Matter* **2024**, *7*, 2382–2398.
- (11) Bailey, M.; Moayedpour, S.; Li, R.; Corrochano-Navarro, A., et al. Deep Batch Active Learning for Drug Discovery. *eLife (reviewed preprint)* **2023**,
- (12) De Ath, G.; Everson, R. M.; Rahat, A. A. M.; Fieldsend, J. E. Greed Is Good: Exploration and Exploitation Trade-offs in Bayesian Optimisation. *ACM Transactions on Evolutionary Learning and Optimization* **2021**, *1*, 1–22.
- (13) Jones, D. R.; Schonlau, M.; Welch, W. J. Efficient Global Optimization of Expensive Black-Box Functions. *Journal of Global Optimization* **1998**, *13*, 455–492.
- (14) Shahriari, B.; Swersky, K.; Wang, Z.; Adams, R. P.; de Freitas, N. Taking the Human Out of the Loop: A Review of Bayesian Optimization. *Proceedings of the IEEE* **2016**, *104*, 148–175.
- (15) Vovk, V.; Gammerman, A.; Shafer, G. *Algorithmic Learning in a Random World*; Springer: New York, 2005.
- (16) Lei, J.; G’Sell, M.; Rinaldo, A.; Tibshirani, R. J.; Wasserman, L. Distribution-Free Predictive Inference for Regression. *Journal of the American Statistical Association* **2018**, *113*, 1094–1111.
- (17) Angelopoulos, A. N.; Bates, S. Conformal Prediction: A Gentle Introduction. *Foundations and Trends® in Machine Learning* **2023**, *16*, 494–591.
- (18) Norinder, U.; Carlsson, L.; Boyer, S.; Eklund, M. Introducing Conformal Prediction in Predictive Modeling. A Transparent and Flexible Alternative to Applicability Domain Determination. *Journal of Chemical Information and Modeling* **2014**, *54*, 1596–1603.
- (19) Arvidsson McShane, S.; Norinder, U.; Alvarsson, J.; Ahlberg, E.; Carlsson, L.;

- Spjuth, O. CPSign: conformal prediction for cheminformatics modeling. *Journal of Cheminformatics* **2024**, *16*.
- (20) Sun, J.; Carlsson, L.; Ahlberg, E.; Norinder, U.; Engkvist, O.; Chen, H. Applying Mondrian Cross-Conformal Prediction To Estimate Prediction Confidence on Large Imbalanced Bioactivity Data Sets. *Journal of Chemical Information and Modeling* **2017**, *57*, 1591–1598.
- (21) Sheridan, R. P. Time-Split Cross-Validation as a Method for Estimating the Goodness of Prospective Prediction. *Journal of Chemical Information and Modeling* **2013**, *53*, 783–790.
- (22) Wu, Z.; Ramsundar, B.; Feinberg, E.; Gomes, J.; Geniesse, C.; Pappu, A. S.; Leswing, K.; Pande, V. MoleculeNet: a benchmark for molecular machine learning. *Chemical Science* **2018**, *9*, 513–530.
- (23) Guo, Q.; Hernandez-Hernandez, S.; Ballester, P. J. Scaffold Splits Overestimate Virtual Screening Performance. Artificial Neural Networks and Machine Learning – ICANN 2024. 2024; pp 58–72.
- (24) Guo, Q.; Hernandez-Hernandez, S.; Ballester, P. J. UMAP-based clustering split for rigorous evaluation of AI models for virtual screening on cancer cell lines. *Journal of Cheminformatics* **2025**, *17*, 39.
- (25) Breiman, L. Random Forests. *Machine Learning* **2001**, *45*, 5–32.
- (26) Pedregosa, F. et al. Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research* **2011**, *12*, 2825–2830.
- (27) Lakshminarayanan, B.; Pritzel, A.; Blundell, C. Simple and Scalable Predictive Uncertainty Estimation using Deep Ensembles. Advances in Neural Information Processing Systems (NeurIPS). 2017.

- (28) Scalia, G.; Grambow, C. A.; Pernici, B.; Li, Y.-P.; Green, W. H. Evaluating Scalable Uncertainty Estimation Methods for Deep Learning-Based Molecular Property Prediction. *Journal of Chemical Information and Modeling* **2020**, *60*, 2697–2717.
- (29) Hirschfeld, L.; Swanson, K.; Yang, K.; Barzilay, R.; Coley, C. W. Uncertainty Quantification Using Neural Networks for Molecular Property Prediction. *Journal of Chemical Information and Modeling* **2020**, *60*, 3770–3780.
- (30) Gal, Y.; Ghahramani, Z. Dropout as a Bayesian Approximation: Representing Model Uncertainty in Deep Learning. Proceedings of the 33rd International Conference on Machine Learning (ICML). 2016; pp 1050–1059.
- (31) Soleimany, A. P.; Amini, A.; Goldman, S.; Rus, D.; Bhatia, S. N.; Coley, C. W. Evidential Deep Learning for Guided Molecular Property Prediction and Discovery. *ACS Central Science* **2021**, *7*, 1356–1367.
- (32) Busk, J.; Jørgensen, P. B.; Bhowmik, A.; Schmidt, M. N.; Winther, O.; Vegge, T. Calibrated uncertainty for molecular property prediction using ensembles of message passing neural networks. *Machine Learning: Science and Technology* **2022**, *3*, 015012.
- (33) Truchon, J.-F.; Bayly, C. I. Evaluating Virtual Screening Methods: Good and Bad Metrics for the “Early Recognition” Problem. *Journal of Chemical Information and Modeling* **2007**, *47*, 488–508.
- (34) Ash, J. R.; Hughes-Oliver, J. M. Confidence bands and hypothesis tests for hit enrichment curves. *Journal of Cheminformatics* **2022**, *14*, 78.
- (35) Janet, J. P.; Duan, C.; Yang, T.; Nandy, A.; Kulik, H. J. A quantitative uncertainty metric controls error in neural network-driven chemical discovery. *Chemical Science* **2019**, *10*, 7913–7922.

- (36) Yang, C.-I.; Li, Y.-P. Explainable uncertainty quantifications for deep learning-based molecular property prediction. *Journal of Cheminformatics* **2023**, *15*.